

tiny Machine Learning



Intelligent edge devices with rich sensors (e.g., billions of mobile phones and IoT devices) have been ubiquitous in our lives.


Deploying artificial intelligence (AI) to these edge devices is attractive for many real-world applications: smart home, smart retail, smart manufacturing, autonomous driving, and more.

However, state-of-the-art deep learning systems typically require tremendous amount of resources (e.g., memory, computation, data, AI experts), both for training and inference.

This hinders the application powerful deep learning systems on edge devices.


TinyML project aims to improve the efficiency of deep learning by requiring ***less computation, fewer engineers,*** and ***less data,*** from algorithm to hardware, from inference to training. Embrace the era of edge AI and AIoT.

Common Challenges



a) Regression

Regression




b) Classification

Classification




Regression

What is the temperature going to be tomorrow?



Classification

Will it be Cold or Hot tomorrow?




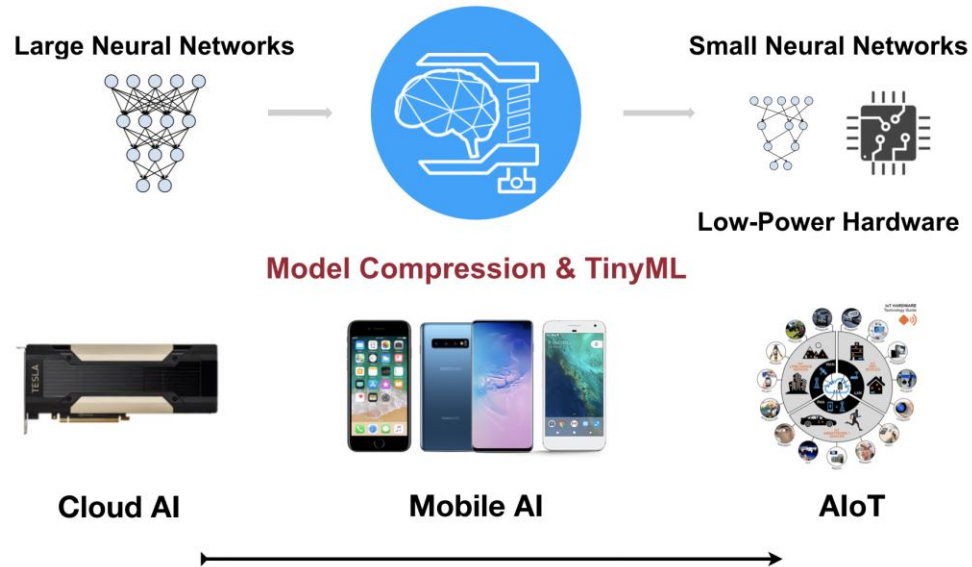

tiny Machine Learning


Machine Learning vs Deep Learning

- DecisionTree
- RandomForest
- XGBoost
- GaussianNB
- Support Vector Machines (SVC and OneClassSVM)
- Relevant Vector Machines (from `skbayes.rvm_ard_models` package)
- SEFR
- PCA

- CNN
- DNN
- RNN
- AutoEncoder
- GAN
- Transformers







Model Compression Techniques

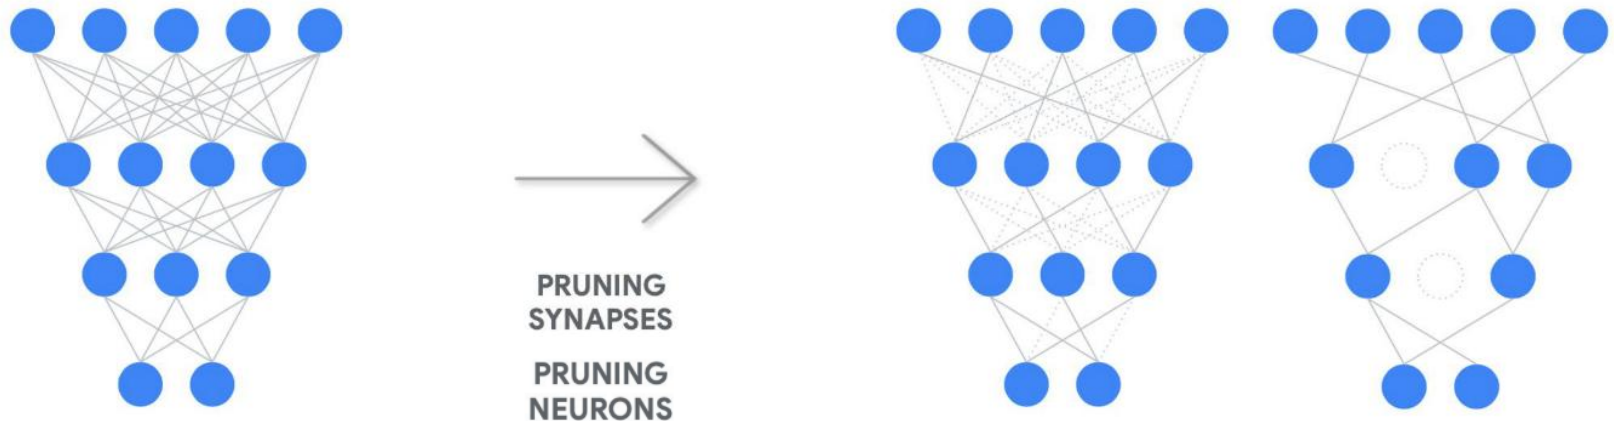
Pruning

Quantization

Knowledge Distillation


...

Pruning



Weight pruning means eliminating unnecessary values in the weight tensors. Practically setting the neural network parameters' values to zero to remove what we estimate are unnecessary connections between the layers of a neural network.

This is done during the training process to allow the neural network to adapt to the changes.



Model Compression Techniques

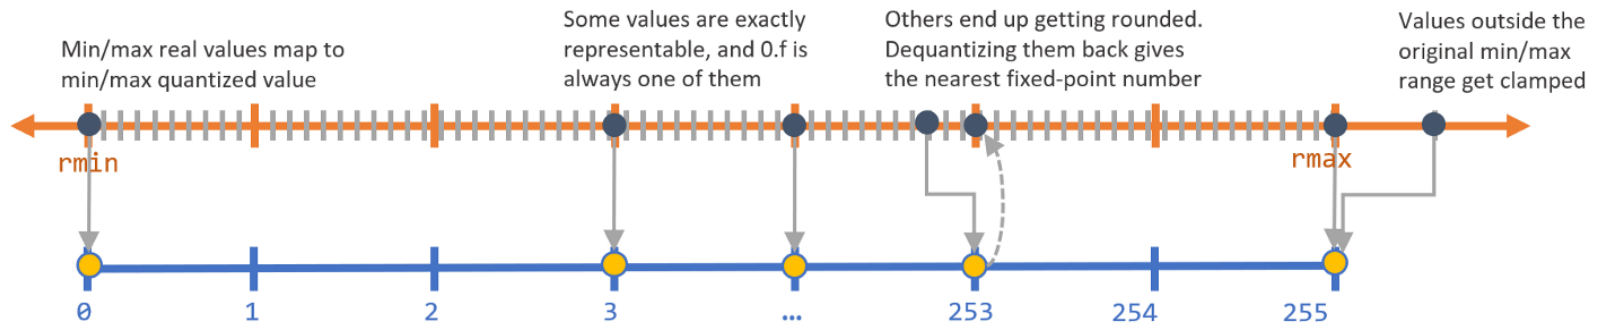
Pruning

Quantization

Knowledge Distillation


...

Quantization



Quantization is the process of transforming an ML model into an equivalent representation that uses parameters and computations at a lower precision. This improves the model's execution performance and efficiency. For example, TensorFlow Lite 8-bit integer quantization results in models that are up to 4x smaller in size, 1.5x-4x faster in computations, and lower power consumption on CPUs.

However, the process of going from higher to lower precision is lossy in nature. As seen in the image below, quantization squeezes a small range of floating-point values into a fixed number of information buckets.



Model Compression Techniques

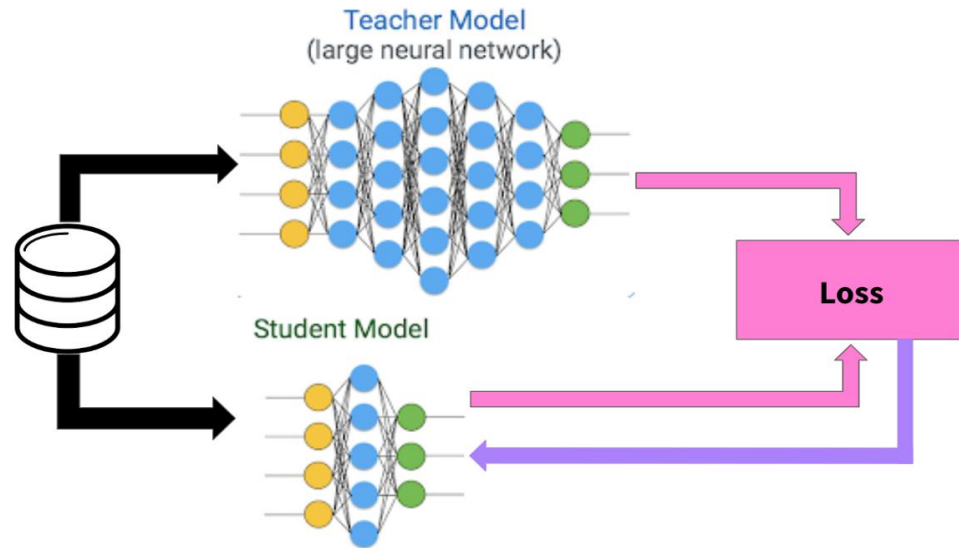
Pruning

Quantization

Knowledge Distillation

...

Knowledge distillation

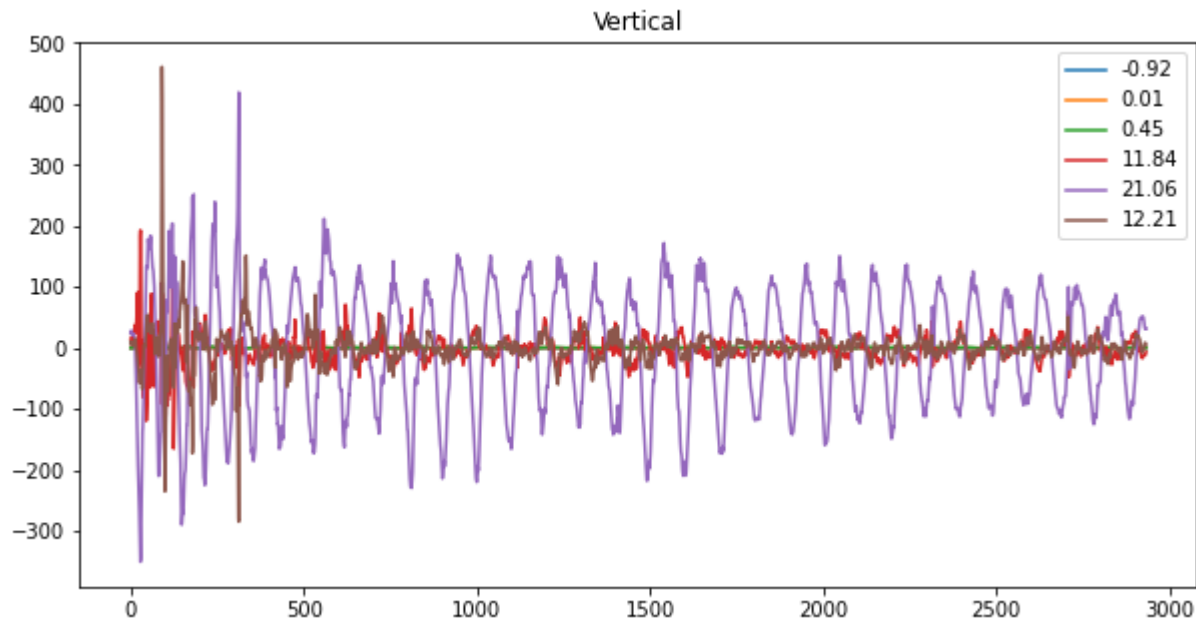


Knowledge Distillation is a procedure for model compression, in which a small (student) model is trained to match a large pre-trained (teacher) model. Knowledge is transferred from the teacher model to the student by minimizing a loss function, aimed at matching softened teacher logits as well as ground-truth labels.

Learn by doing...

Machine Learning version vs Deep Learning Model

Gesture Recognition: the Machine Learning way



DEMO

```
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from everywhere.ml.sklearn.ensemble import RandomForestClassifier

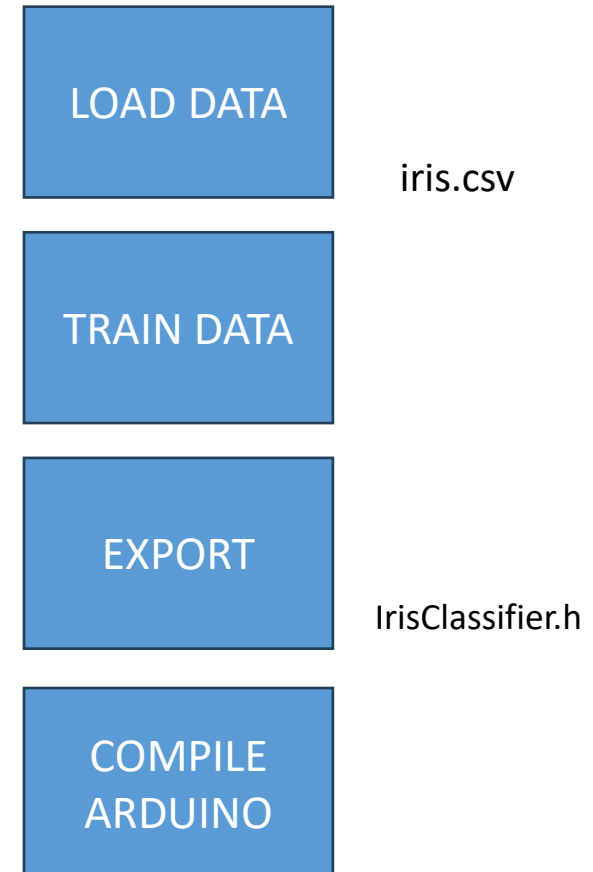
def load_data_from_csv(filename: str, label_column: str) -> tuple:
    """
    Convert csv file to X and y
    :param label_column:
    :param filename:
    :return:
    """
    df = pd.read_csv(filename)
    x_columns = [c for c in df.columns if c != label_column]
    X = df[x_columns].to_numpy(dtype=float)
    y_string = df[label_column]
    label_encoder = LabelEncoder().fit(y_string)
    y_numeric = label_encoder.transform(y_string)
    print('Label mapping', {label: i for i, label in enumerate(label_encoder.classes_)})

    return X, y_numeric

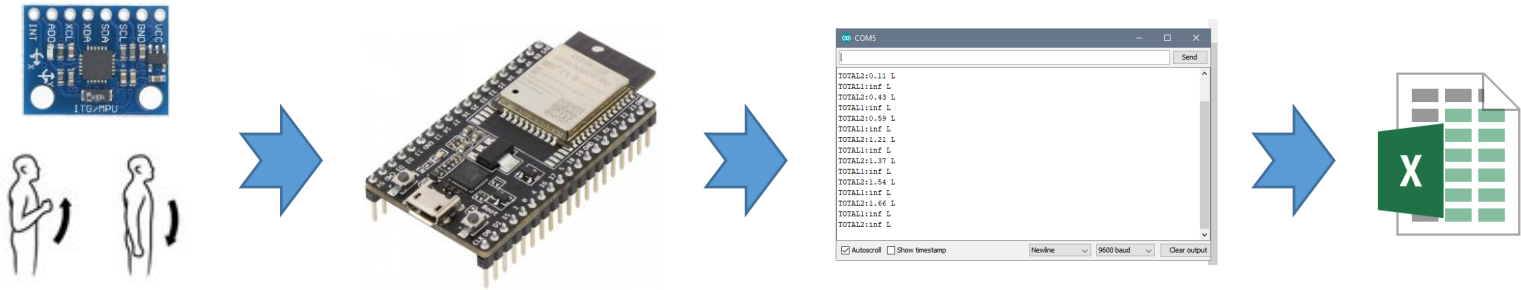
X, y = load_data_from_csv('iris.csv', label_column='variety')
clf = RandomForestClassifier(n_estimators=5).fit(X, y)

import sys
sys.stdout = open('IrisClassifier.h', 'wt')
print(clf.to_arduino(instance_name='irisClassifier'))
```

pip install everywhere.ml



Exercise: Capture and train data



Step 1.1 Record the "no gesture" gesture
Arm at rest position
Move slowly to add some noise
min. of 20 secs
save data from terminal to file "rest.csv"

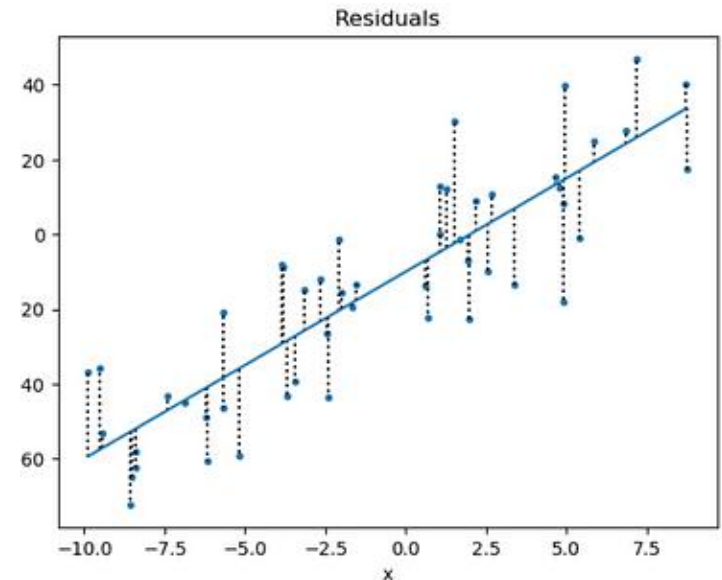
Record.ino

Step 1.2 Record all the other gestures
4 different movements (Left, Right, Front, Up)
min. of 20 secs for each motion/repetition
save data from terminal to file "xxx.csv"

tiny Machine Learning

Linear Regression


...is a statistical method used to model the relationship between a dependent variable and one or more independent variables by fitting a linear equation to observed data. It is one of the simplest and most used techniques in statistical modelling and machine learning for predictive analysis.



tiny Machine Learning

Linear Regression

...relies on several key assumptions to provide reliable and valid results.




A. **Linearity:** The relationship between dependent and independent variables should follow a straight-line pattern, meaning changes in the independent variables proportionally affect the dependent variable.

B. **Independence of Errors:** Residuals should not be correlated with each other; each observation's error should be independent. Autocorrelation may indicate missing variables or model issues.

tiny Machine Learning

Linear Regression



C. **Homoscedasticity:** The variance of errors should remain constant across all values of the independent variables. Inconsistent variance (heteroscedasticity) can lead to biased estimates.


D. **Normality of Errors:** Residuals should follow a normal distribution to ensure accurate statistical inference, hypothesis testing, and confidence intervals.

E. **No Perfect Multicollinearity:** Independent variables should not be perfectly correlated, as this makes coefficient estimation unstable and unreliable.

F. **Zero Conditional Mean:** The average error should be zero across all values of independent variables, ensuring unbiased predictions without systematic over- or underestimation.

tiny Machine Learning

Linear Regression - Model




$$Y = \beta_0 + \beta_1 X + \varepsilon$$

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n + \varepsilon$$

- Y is the dependent variable.
- X is the independent variable.
- β_0 is the intercept (the value of Y when $X=0$).
- β_1 is the slope (the change in Y for a one-unit change in X).
- ε represents the error term, which captures the difference between the observed and predicted values.


tiny Machine Learning

Linear Regression - Coefficients



The coefficients $\beta_0, \beta_1, \dots, \beta_n$ are estimated using the method of least squares. This method minimizes the sum of the squared differences between the observed values of the dependent variable and the values predicted by the linear model.

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$





- y_i is the observed value of the dependent variable for the i th observation.
- \hat{y}_i is the predicted value of the dependent variable for the i th observation based on the linear regression model.
- n is the number of observations.

tiny Machine Learning

Linear Regression – Minimizing the Objective Function

Gradient descent is an iterative optimization algorithm used to minimize the MSE by adjusting the coefficients in the direction of the negative gradient of the MSE with respect to each coefficient.



$$\beta_j^{(t+1)} = \beta_j^{(t)} - \alpha \frac{\partial}{\partial \beta_j} \text{MSE}$$


tiny Machine Learning

Linear Regression -

In practice (step by step)



Importing libraries
Load Dataset
Exploratory Data Analysis
Splitting the data
Feature selection
Training the model
Evaluating the model
Export Model
Run in Arduino




tiny Machine Learning

Linear Regression -

In practice (step by step)

```
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.calibration import LabelEncoder
from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import RFE
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from micromlgen import port
import pylab
import scipy.stats as stats
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```



tiny Machine Learning

Load dataset

```
data = load_diabetes() # load data
```

Create a DataFrame

```
df_diabetes = pd.DataFrame(data.data, columns=data.feature_names)
```

Add target variable to the DataFrame

```
df_diabetes['target'] = data.target
```

Remove NaN values


```
df = df_diabetes.dropna(axis='rows') #remove NaN
```

Display the DataFrame

```
df_diabetes.head()
```


```
df_diabetes.info()
```

```
df_diabetes.describe()
```



tiny Machine Learning

	age	sex	bmi	bp	s1	s2	s3	s4	s5	s6	target
count	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	4.420000e+02	442.000000
mean	-2.511817e-19	1.230790e-17	-2.245564e-16	-4.797570e-17	-1.381499e-17	3.918434e-17	-5.777179e-18	-9.042540e-18	9.293722e-17	1.130318e-17	152.133484
std	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	4.761905e-02	77.093005
min	-1.072256e-01	-4.464164e-02	-9.027530e-02	-1.123988e-01	-1.267807e-01	-1.156131e-01	-1.023071e-01	-7.639450e-02	-1.260971e-01	-1.377672e-01	25.000000
25%	-3.729927e-02	-4.464164e-02	-3.422907e-02	-3.665608e-02	-3.424784e-02	-3.035840e-02	-3.511716e-02	-3.949338e-02	-3.324559e-02	-3.317903e-02	87.000000
50%	5.383060e-03	-4.464164e-02	-7.283766e-03	-5.670422e-03	-4.320866e-03	-3.819065e-03	-6.584468e-03	-2.592262e-03	-1.947171e-03	-1.077698e-03	140.500000
75%	3.807591e-02	5.068012e-02	3.124802e-02	3.564379e-02	2.835801e-02	2.984439e-02	2.931150e-02	3.430886e-02	3.243232e-02	2.791705e-02	211.500000
max	1.107267e-01	5.068012e-02	1.705552e-01	1.320436e-01	1.539137e-01	1.987880e-01	1.811791e-01	1.852344e-01	1.335973e-01	1.356118e-01	346.000000



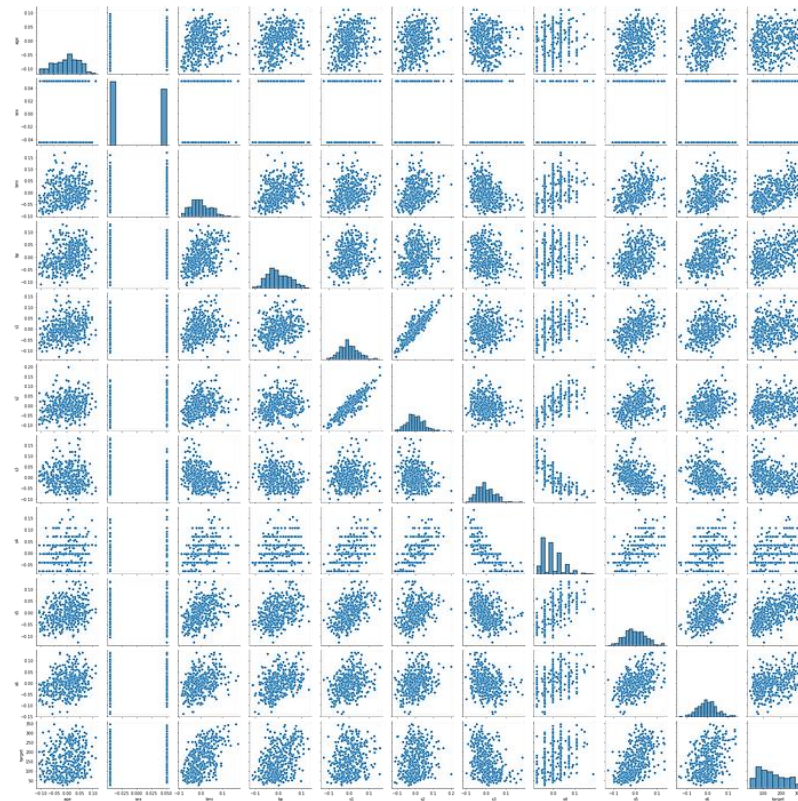

	age	sex	bmi	bp	s1	s2	s3	s4	s5	s6	target
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	-0.002592	0.019907	-0.017646	151.0
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	-0.039493	-0.068332	-0.092204	75.0
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	-0.002592	0.002861	-0.025930	141.0
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	0.034309	0.022688	-0.009362	206.0
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	-0.002592	-0.031988	-0.046641	135.0

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 442 entries, 0 to 441
Data columns (total 11 columns):
 #   Column      Non-Null Count  Dtype
---  ---
 0   age         442 non-null    float64
 1   sex         442 non-null    float64
 2   bmi         442 non-null    float64
 3   bp          442 non-null    float64
 4   s1          442 non-null    float64
 5   s2          442 non-null    float64
 6   s3          442 non-null    float64
 7   s4          442 non-null    float64
 8   s5          442 non-null    float64
 9   s6          442 non-null    float64
10  target      442 non-null    float64
dtypes: float64(11)
memory usage: 38.1 KB
```


tiny Machine Learning

Seaborn analysis

```
sns.pairplot(df_diabetes)
```



tiny Machine Learning

```
df = df_diabetes.iloc[:100,0:10]
```

```
X=df.to_numpy()
```


```
# Converting string value to int type for labels: Setosa = 0, Versicolor = 1
```

```
y=df_diabetes.iloc[:100,-1]
```

```
y = LabelEncoder().fit_transform(y)
```

```
# Split the data into training and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
random_state=42)
```



tiny Machine Learning

```
df = df_diabetes.iloc[:100,0:10]
```

```
X=df.to_numpy()
```


```
# Converting string value to int type for labels: Setosa = 0, Versicolor = 1
```

```
y=df_diabetes.iloc[:100,-1]
```


```
y = LabelEncoder().fit_transform(y)
```

```
# Split the data into training and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
random_state=42)
```



tiny Machine Learning



Recursive Feature Elimination (RFE)

Build a logreg and compute the feature importances

```
linearRegressor = LinearRegression()
```

create the RFE model and select 5 attributes

```
rfe = RFE(estimator=linearRegressor, n_features_to_select=5, step=30)
```

```
rfe = rfe.fit(X_train, y_train)
```

summarize the selection of the attributes

```
print('Selected features: %s' % list(df.columns[rfe.support_]))
```

```
selected_features = ['sex', 'bmi', 's1', 's5']
```



```
df_selected = df[selected_features]
```

```
plt.subplots(figsize=(8, 5))
```


```
sns.heatmap(df_selected.corr(), annot=True, cmap="RdYlGn")
```

```
plt.show()
```

tiny Machine Learning



tiny Machine Learning



Ordinary Least Squares (OLS)

... regression analysis is a statistical method used to estimate the relationship between one or more independent variables (predictors) and a dependent variable (response).

```
# Fit the regression line using 'OLS'
```

```
lr = sm.OLS(y_train, X_train).fit()
```

```
# Add a constant to get an intercept
```

```
X_train_sm = sm.add_constant(X_train)
```

```
# Fit the regression line using 'OLS'
```

```
lr = sm.OLS(y_train, X_train_sm).fit()
```

```
# Print the parameters, i.e. the intercept and the slope of the regression line fitted
```


```
parameters = lr.params
```

```
print('Model fit parameter: ', parameters)
```

```
# Performing a summary operation lists out all the different parameters
```

```
print(lr.summary())
```

tiny Machine Learning



Ordinary Least Squares (OLS)

... regression analysis is a statistical method used to estimate the relationship between one or more independent variables (predictors) and a dependent variable (response).



```
predict = parameters[0] +  
           parameters[0]*X_train[:,0]+  
           parameters[1]*X_train[:,1]+  
           parameters[2]*X_train[:,2]+  
           parameters[3]*X_train[:,3]+  
           parameters[4]*X_train[:,4]+  
           parameters[5]*X_train[:,5]+  
           parameters[6]*X_train[:,6]+  
           parameters[7]*X_train[:,7]+  
           parameters[8]*X_train[:,8]+  
           parameters[9]*X_train[:,9]
```

```
plt.plot(y_train)  
plt.plot(predict, 'r')  
plt.grid()  
plt.show()
```


tiny Machine Learning

Ordinary Least Squares (OLS)

... regression analysis is a statistical method used to estimate the relationship between one or more independent variables (predictors) and a dependent variable (response).




tiny Machine Learning



Ordinary Least Squares (OLS)

... regression analysis is a statistical method used to estimate the relationship between one or more independent variables (predictors) and a dependent variable (response).

```
residuals = y_train-predict  
sns.distplot(residuals)  
plt.grid()
```





tiny Machine Learning

Ordinary Least Squares (OLS)

... regression analysis is a statistical method used to estimate the relationship between one or more independent variables (predictors) and a dependent variable (response).

```
stats.probplot(residuals, dist="norm", plot=pylab)
plt.grid()
pylab.show()
```





tiny Machine Learning

Model Evaluation

```
model = linearRegressor.fit(X_train,y_train)
```

```
predict = model.predict(X_train)
```

```
x_ax = range(len(y_train))  
plt.plot(x_ax, y_train, label="original")  
plt.plot(x_ax, predict, label="predicted")  
plt.xlabel('X-axis')  
plt.ylabel('Y-axis')  
plt.legend(loc='best',fancybox=True, shadow=True)  
plt.grid(True)  
plt.show()
```




tiny Machine Learning

Model Evaluation

```
score = model.score(X_train, y_train)
training_predict = model.predict(X_train)
mse = mean_squared_error(y_train, training_predict)
```

```
print("R-squared:", score)
print("MSE: ", mse)
print("RMSE: ", mse*(1/2.0))
```



```
R-squared: 0.4665605455873243
MSE: 261.72727823900914
RMSE: 130.86363911950457
```



tiny Machine Learning

Residuals



```
residuals = y_train-training_predict  
sns.distplot(residuals)  
plt.grid()
```



tiny Machine Learning

Residuals



```
stats.probplot(residuals, dist="norm", plot=pylab)  
plt.grid()  
pylab.show()
```




tiny Machine Learning

Residuals

```
x_ax = range(len(y_test))
plt.plot(x_ax, y_test, label="original")
plt.plot(x_ax, test_predict, label="predicted")
plt.xlabel('X-axis')
plt.ylabel('Y-axis')
plt.legend(loc='best', fancybox=True, shadow=True)
plt.grid(True)
plt.show()
```



tiny Machine Learning



Residuals

```
score = model.score(X_test, y_test)
test_predict = model.predict(X_test)
mse = mean_squared_error(y_test, test_predict)
```



```
print("R-squared:", score)
print("MSE: ", mse)
print("RMSE: ", mse*(1/2.0))
```

```
R-squared: 0.267072213466814
MSE: 300.75203101864935
RMSE: 150.37601550932467
```


tiny Machine Learning

Residuals



```
residuals = y_test-test_predict  
sns.distplot(residuals)  
plt.grid()
```



tiny Machine Learning

Residuals

```
stats.probplot(residuals, dist="norm", plot=pylab)  
plt.grid()  
pylab.show()
```



tiny Machine Learning

Export & Arduino


```
print(port(model))
```

```
with open('.\LinearRegressor.h', 'w') as file:  
    file.write(port(model))
```


NOTE: in case of Arduino Compile Errors

In file LinearRegressor.h verify if include is :

```
#include <cstdarg>    => change to => #include <stdarg.h>
```



tiny Machine Learning



Export & Arduino

```
#include "LinearRegressor.h"
```

```
Eloquent::ML::Port::LinearRegression LinearRegressor;
```

```
void setup()
```

```
{
```

```
Serial.begin(115200);
```

```
}
```

```
void loop()
```

```
{
```

```
float X_1[] = {0.02717829, 0.05068012, 0.01750591, -0.03321323, -0.00707277, 0.04597154,  
-0.06549067, 0.07120998, -0.09643495, -0.05906719};
```

```
int result_1 = LinearRegressor.predict(X_1);
```

```
Serial.print("Result of predict with input X1 (real value = 13):");
```

```
Serial.println(result_1);
```

```
delay(2000);
```

```
}
```



tiny Machine Learning CNN

tiny Machine Learning



```
pip install "everywhereml>=0.2.32"  
pip install tensorflow
```

```
import numpy as np
```

```
from sklearn.datasets import load_iris  
from sklearn.preprocessing import MinMaxScaler
```

```
import tensorflow as tf  
from tensorflow.keras import layers  
import matplotlib.pyplot as plt
```

```
from everywhereml.code_generators.tensorflow import convert_model
```

tiny Machine Learning



```
def load_Xy():  
    """  
    Load data  
    It may be a CSV file or whatever  
    """  
    return load_iris(return_X_y=True)
```

```
def get_Xy():  
    """  
    Normalize X  
    One-hot encode y  
    """  
    X, y = load_Xy()  
    X = np.asarray(X)  
    y = np.asarray(y)  
    num_classes = y.max()  
    eye = np.eye(num_classes + 1)  
    X_norm = MinMaxScaler().fit_transform(X)  
    y_hot = np.asarray([eye[yi] for yi in y], dtype=int)  
  
    return X_norm, y_hot
```

tiny Machine Learning



```
"""
Instantiate NN for classification
"""
def instantiate_nn_for_classification(input_shape, num_classes):
    model = tf.keras.Sequential()
    model.add(layers.Dense(32, activation='relu', input_shape=input_shape))
    model.add(layers.Dense(16, activation='relu'))
    model.add(layers.Dense(num_classes, activation='softmax'))
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

    return model
```

```
"""
Split data between train, validation and test
"""
from sklearn.model_selection import train_test_split

X, y = get_Xy()
X_train, X_test, y_train, y_test = train_test_split(X, y,
    test_size=0.3)
X_train, X_val, y_train, y_val = train_test_split(X_train,
    y_train, test_size=0.3)
```


tiny Machine Learning



```
"""
Instantiate NN for classification
"""
def instantiate_nn_for_classification(input_shape, num_classes):
    model = tf.keras.Sequential()
    model.add(layers.Dense(32, activation='relu', input_shape=input_shape))
    model.add(layers.Dense(16, activation='relu'))
    model.add(layers.Dense(num_classes, activation='softmax'))
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

    return model
```

```
"""
Split data between train, validation and test
"""
from sklearn.model_selection import train_test_split

X, y = get_Xy()
X_train, X_test, y_train, y_test = train_test_split(X, y,
    test_size=0.3)
X_train, X_val, y_train, y_val = train_test_split(X_train,
    y_train, test_size=0.3)
```

tiny Machine Learning



```
"""
Train model
"""
input_shape = X.shape[1:]
num_classes = y.shape[1]

model = instantiate_nn_for_classification(input_shape, num_classes)
model.summary()

history = model.fit(X_train, y_train, epochs=50, batch_size=16, validation_data=(X_val, y_val))

print(history.history.keys())

# plot history accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

tiny Machine Learning



```
"""
```

```
Export NN to C++
```

```
Copy-paste the generated code inside a file named model.h or irisModel.h  
in your Arduino project
```

```
"""
```

```
from everywhere.ml.code_generators.tensorflow import convert_model
```

```
out = convert_model(model, X, y, model_name='irisModel')  
print(out)
```

```
with open('.\cnn.h', 'w') as file:  
    file.write(out)
```

tiny Machine Learning

ARDUINO

install libs =>



EloquentTinyML by Simone Salerno, eloquentarduino@gmail.com
3.0.1 installed
An eloquent interface to Tensorflow Lite for Microcontrollers
[More info](#)
3.0.1 **REMOVE**

tflm_esp32 by Simone Salerno, eloquentarduino@gmail.com
2.0.0 installed
TensorFlow for ESP32
[More info](#)
2.0.0 **REMOVE**

```
#include "cnn.h"
#include <tflm_esp32.h>
#include <eloquent_tinyml.h>

#define ARENA_SIZE 2000 // this is trial-and-error process
Eloquent::TF::Sequential<TF_NUM_OPS, ARENA_SIZE> tf;
...
```

tiny Machine Learning



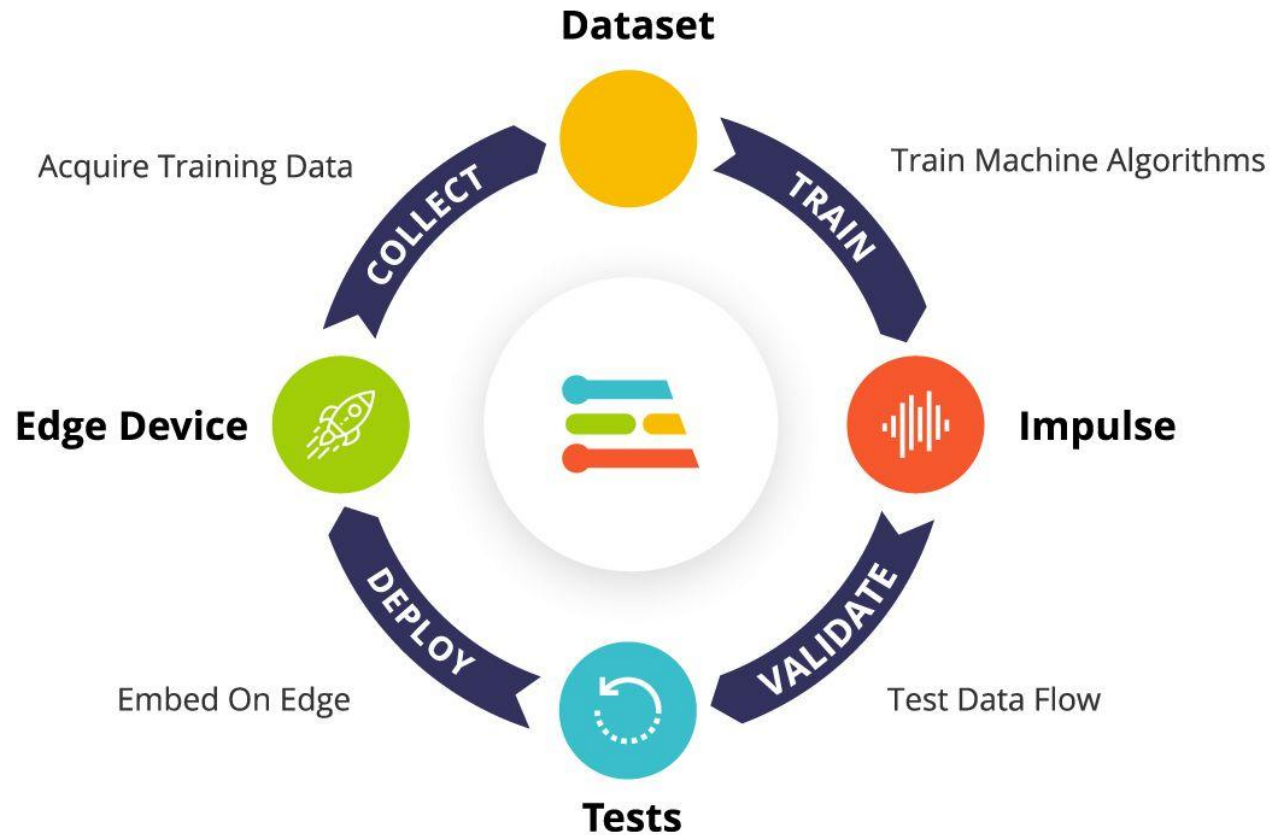
```
void setup() {  
    Serial.begin(115200);  
    delay(3000);  
    Serial.println("__TensorFlow IRIS__");  
  
    // configure input/output  
    tf.setNumInputs(4);  
    tf.setNumOutputs(3);  
  
    // add required ops  
    tf.resolver.AddFullyConnected();  
    tf.resolver.AddSoftmax();  
  
    while (!tf.begin(irisModel).isOk())  
        Serial.println(tf.exception.toString());  
}
```

tiny Machine Learning

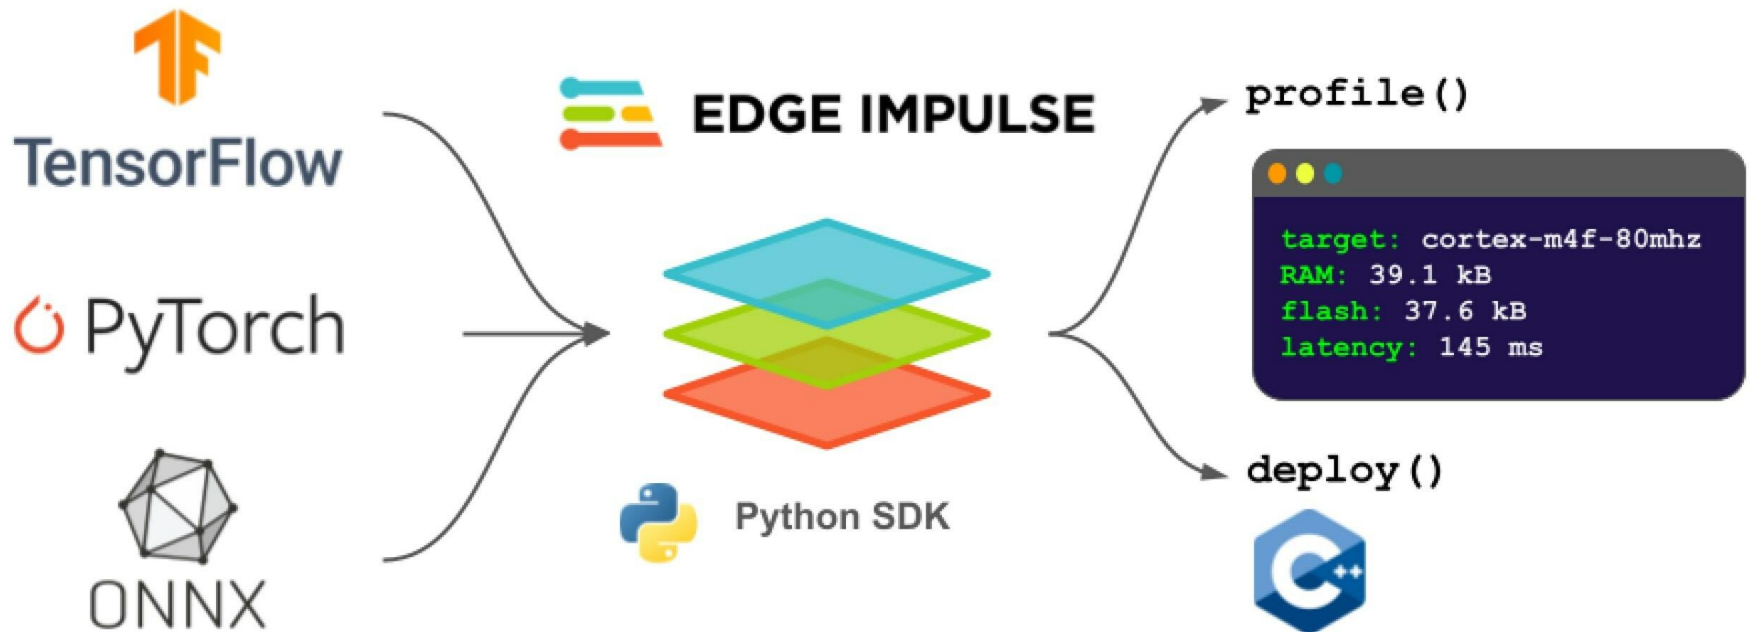


```
void loop() {  
    // classify class 0  
    if (!tf.predict(x0).isOk()) {  
        Serial.println(tf.exception.toString());  
        return;  
    }  
  
    Serial.print("expected class 0, predicted class ");  
    Serial.println(tf.classification);  
  
    // how long does it take to run a single prediction?  
    Serial.print(tf.benchmark.microseconds());  
    Serial.println("us");  
  
    delay(1000);  
}
```

tiny Machine Learning



tiny Machine Learning



Todo: Register and try online...



tiny Machine Learning LSTM

tiny Machine Learning



```
pip install embedded_window  
pip install eloquent_tensorflow
```

```
import numpy as np  
import pandas as pd
```

```
from embedded_window import Window
```

```
from sklearn.model_selection import train_test_split
```

```
import tensorflow as tf  
from tensorflow.keras import layers
```

tiny Machine Learning



```
window_size = 150
window_shift = 20
window = Window(length=window_size,
shift=window_shift)
```

```
def read_file(filename, plot: bool = True):
    df = pd.read_csv(filename)
    acc = df[['ax(mg)', 'ay(mg)', 'az(mg)']]

    if plot:
        acc.plot(title=filename)

    # normalize data
    # in this case, accel is in the range -3, 3
    accel = np.clip(acc.to_numpy() / 3, -1, 1)
    y = np.zeros(len(accel))

    # apply window
    # returns a (N, window_length, accel.columns) array
    accel_windows, _, _ = window.fit_transform(accel, y)
    return accel_windows
```

tiny Machine Learning



```
def get_Xy():  
    """  
    Read files  
    One-hot encode y  
    """  
  
    x_idle = read_file('idle.csv')  
    x_horizontal = read_file('left-right.csv')  
    x_vertical = read_file('up-down.csv')  
  
    y_idle = [0] * len(x_idle)  
    y_horizontal = [1] * len(x_horizontal)  
    y_vertical = [2] * len(x_vertical)  
  
    X = np.vstack((x_idle, x_horizontal, x_vertical))  
    y = np.concatenate((y_idle, y_horizontal, y_vertical))  
    num_classes = y.max()  
    eye = np.eye(num_classes + 1)  
    y_hot = np.asarray([eye[yi] for yi in y], dtype=int)  
  
    return X, y_hot
```

tiny Machine Learning






```
X, y = get_Xy()
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.3)
num_classes = y.shape[1]
print("X.shape", X.shape)
```

```
"""
Print samples of each class
"""

def print_sample(variable_name: str, filename: str):
    sample = read_file(filename, plot=False)[0].flatten()
    values = ', '.join('%0.4f' % x for x in sample)
    print(f'float {variable_name}[{len(sample)}] = {{ {values} }};')

print_sample('idle', 'idle.csv')
print_sample('horizontal', 'left-right.csv')
print_sample('vertical', 'up-down.csv')
```

tiny Machine Learning



tiny Machine Learning



```
model = instantiate_lstm_for_classification(X.shape[1:], num_classes)
print(model.summary())
```

```
"""
Train model
"""
history = model.fit(X_train, y_train, epochs=30, validation_data=(X_val, y_val))

print(history.history.keys())
# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

tiny Machine Learning



```
"""
Evaluate on test data
"""
model.evaluate(X_test, y_test)

"""
Save model
"""
from eloquent_tensorflow import convert_model
out = convert_model(model)
print(out)

with open('.\lstm.h', 'w') as file:
    file.write(out)
```


tiny Machine Learning

ARDUINO

install libs =>




EloquentTinyML by Simone Salerno, eloquentarduino@gmail.com
3.0.1 installed
An eloquent interface to Tensorflow Lite for Microcontrollers
[More info](#)
3.0.1 **REMOVE**

tflm_esp32 by Simone Salerno, eloquentarduino@gmail.com
2.0.0 installed
TensorFlow for ESP32
[More info](#)
2.0.0 **REMOVE**

```
#include <Arduino.h>
#include "lstm.h"
#include <tflm_esp32.h>
#include <eloquent_tinym1.h>
```

tiny Machine Learning



```
// this is trial-and-error process
#define ARENA_SIZE 30000

Eloquent::TF::Sequential<TF_NUM_OPS, ARENA_SIZE> tf;

float idle[450] = { -0.3334, 0.0289, ...
float horizontal[450] = { -0.4638, ...
float vertical[450] = { -0.5387, ...

void predictSample(const char *classLabel, float *input, uint8_t expectedOutput)
{
    // classify class 0
    if (!tf.predict(input).isOk()) {
        Serial.println(tf.exception.toString());
        return;
    }

    Serial.print("Predicting sample of ");
    Serial.print(classLabel);
    Serial.print(": expcted id=");
    Serial.print(expectedOutput);
    Serial.print(", predicted=");
    Serial.println(tf.classification);
}
```

tiny Machine Learning



```
void setup() {  
  Serial.begin(115200);  
  delay(3000);  
  Serial.println("__LSTM__");  
  
  // configure input/output  
  tf.setNumInputs(TF_NUM_INPUTS);  
  tf.setNumOutputs(TF_NUM_OUTPUTS);  
  
  registerNetworkOps(tf);  
  
  while (!tf.begin(tfModel).isOk()) {  
    Serial.println(tf.exception.toString());  
    delay(1000);  
  }  
}
```

tiny Machine Learning




```
void loop() {  
    predictSample("idle", idle, 0);  
    predictSample("horizontal", horizontal, 1);  
    predictSample("vertical", vertical, 2);  
  
    // how long does it take to run a single prediction?  
    Serial.print(tf.benchmark.microseconds());  
    Serial.println("us");  
  
    delay(1000);  
}
```

Final Project



TOPICS

1. Gesture Recognition
2. Person Fall detection
3. Surface recognition
4. Machine Anomaly
5. ...
6. (open idea)



“a thing to save the day...”

Final Project

Define:
goals,
limitations,
challenges.

1. Define the Problem

Identify the problem to solve (e.g., anomaly detection, voice recognition, gesture detection).
Consider constraints like power consumption, latency, and hardware limitations.

2. Select Hardware & Tools

Choose a microcontroller (e.g., Arduino, Raspberry Pi Pico, ESP32).
Select sensors (e.g., accelerometers, microphones, sensors).
Development tools (e.g., TensorFlow Lite for Microcontrollers, Edge Impulse, Arduino IDE).

3. Data Collection & Preprocessing

Gather relevant data (sensor readings, audio, analog/digital).
Clean and normalize data (remove noise, scale values).
Augment data (if needed) (e.g., data balancing, synthetic data).

4. Model Selection & Training

Choose an appropriate ML model (CNN, RNN, ML models).
Train the model.

5. Deploying the Model

Convert the trained model to C/C++.
Deploy to the microcontroller.
Optimize for real-time performance and memory efficiency.

6. Testing & Validation

Run real-time tests on embedded hardware.
Validate performance metrics (accuracy, inference speed, power consumption).